

# CSE419 – Artificial Intelligence and Machine Learning 2020

PhD Furkan Gözükkara, Toros University








<https://github.com/FurkanGozukara/CSE419-Artificial-Intelligence-and-Machine-Learning-2020>

## Lecture 9

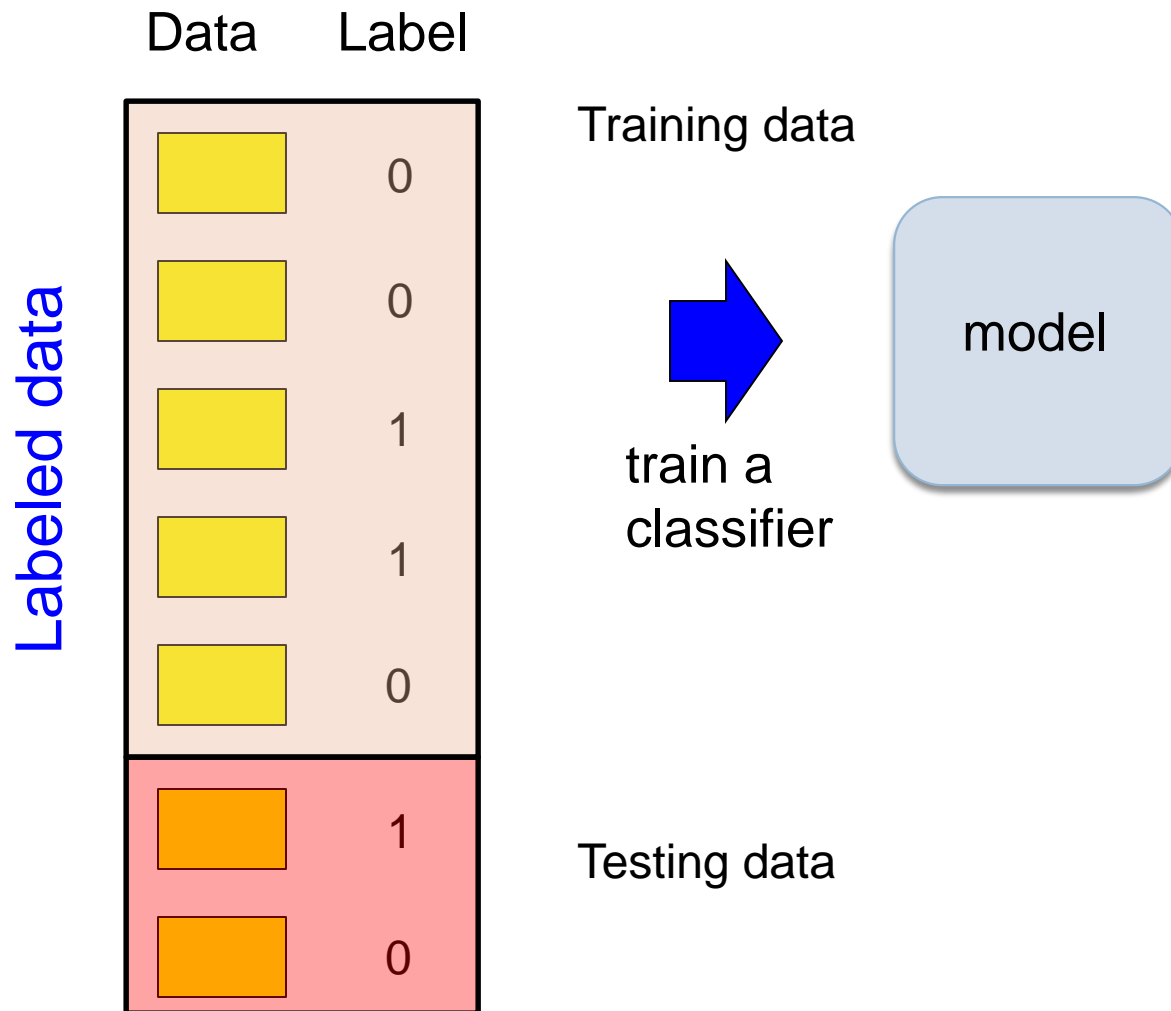
## Evaluation

*Based on Asst. Prof. Dr. David Kauchak (Pomona College) Lecture Slides*

# Supervised evaluation


		Data	Label	
Labeled data				Training data
			0	
			0	
			1	
			1	
			0	
				Testing data
			1	
			0	

# Supervised evaluation



# Supervised evaluation







Data      Label



Pretend like we  
don't know the  
labels

# Supervised evaluation

Data      Label



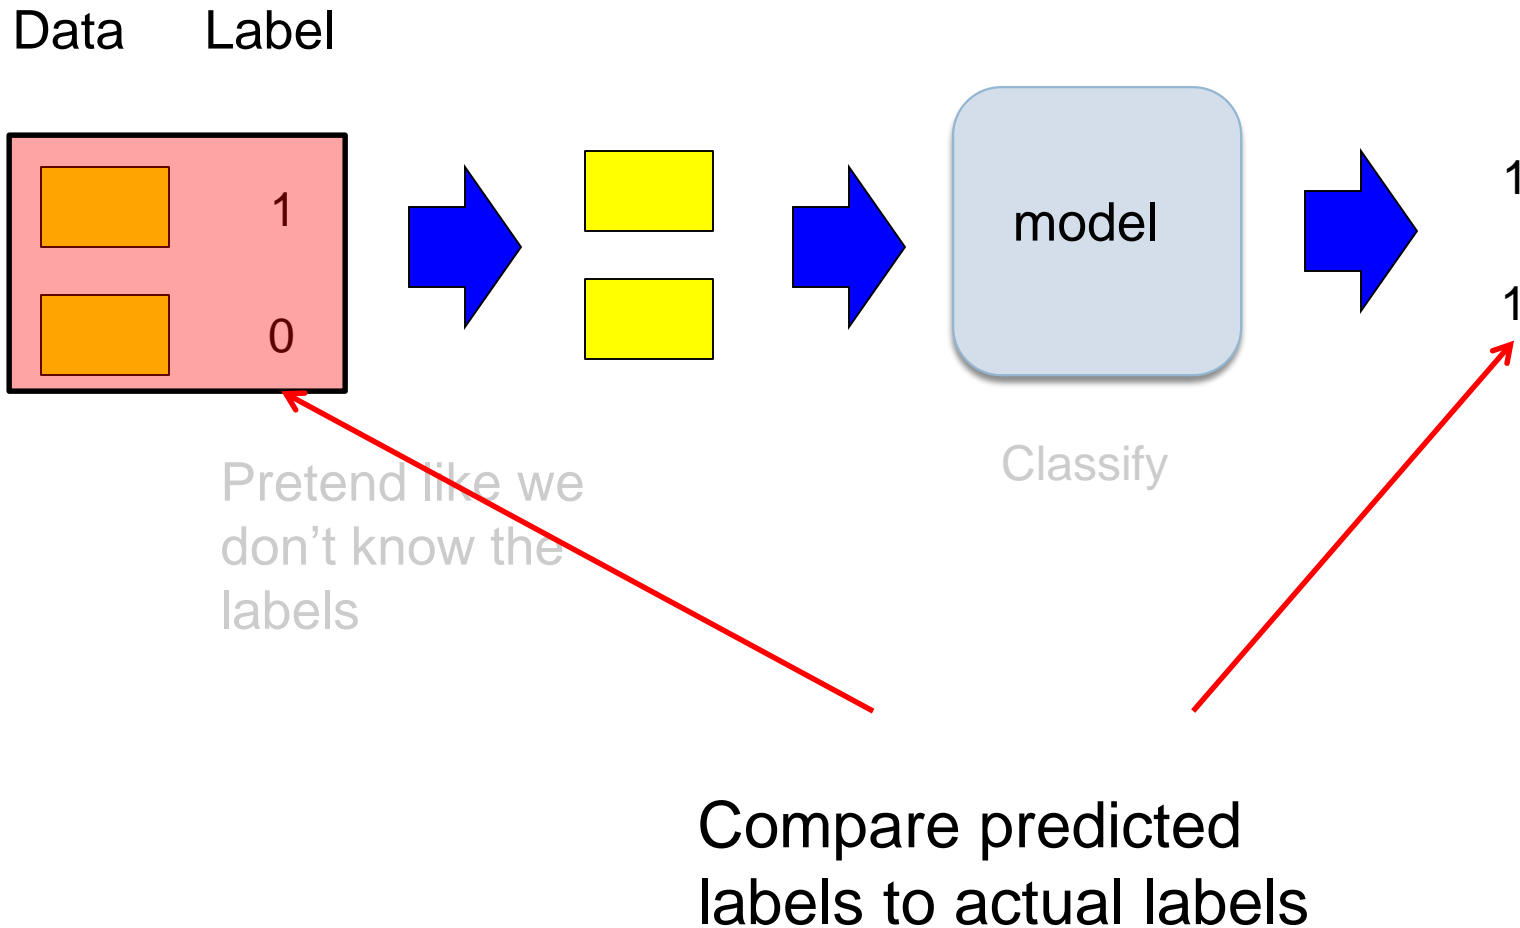
1

1


Pretend like we  
don't know the  
labels

Classify

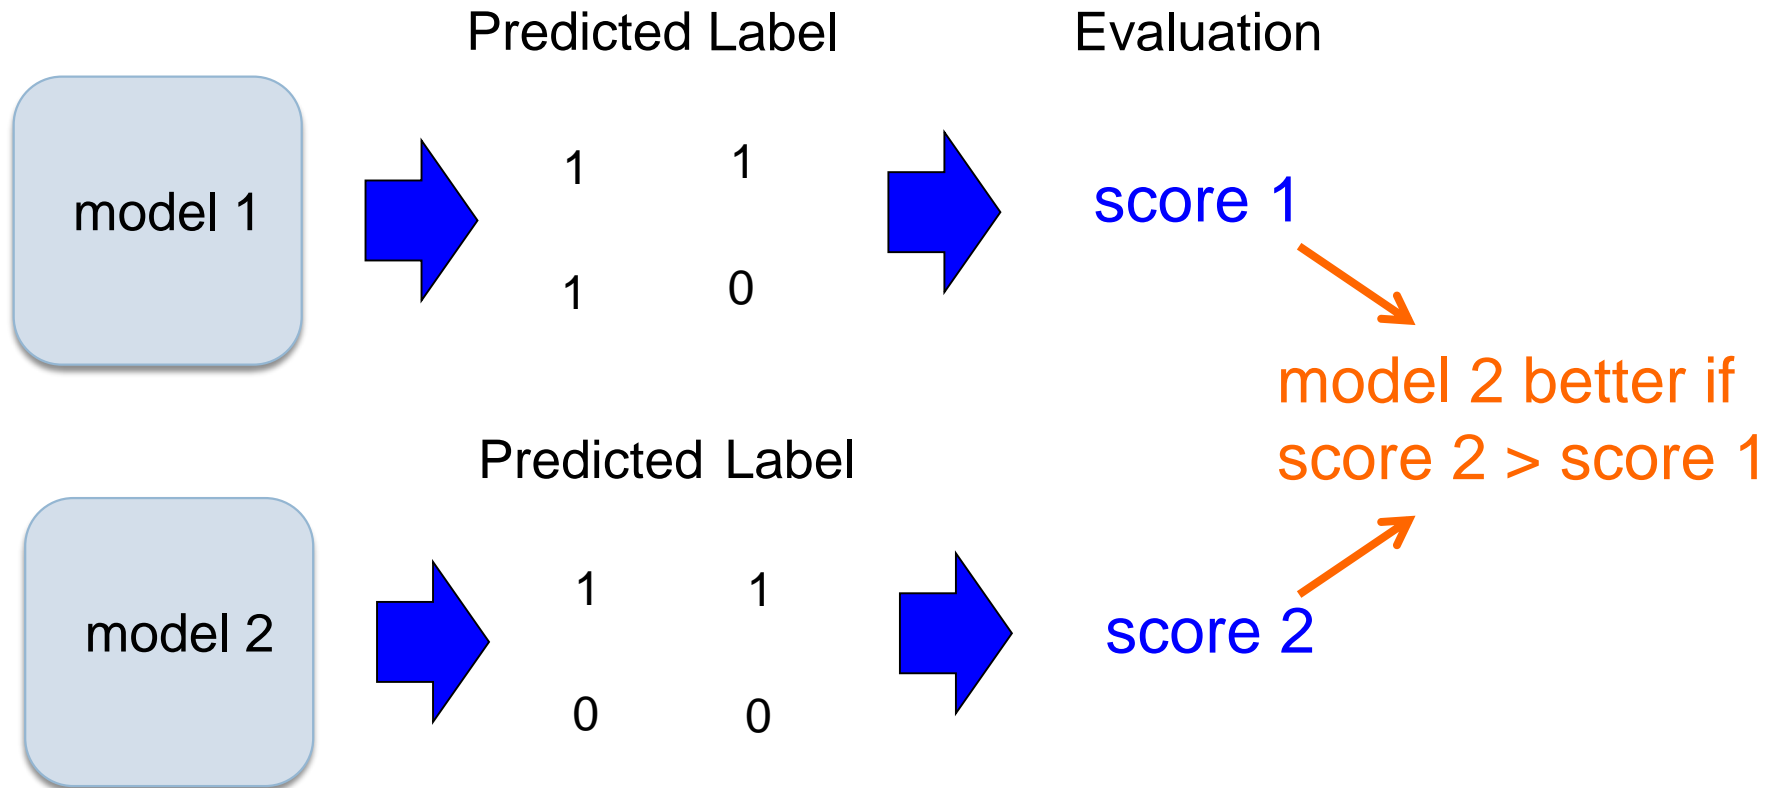
# Supervised evaluation



# Comparing algorithms



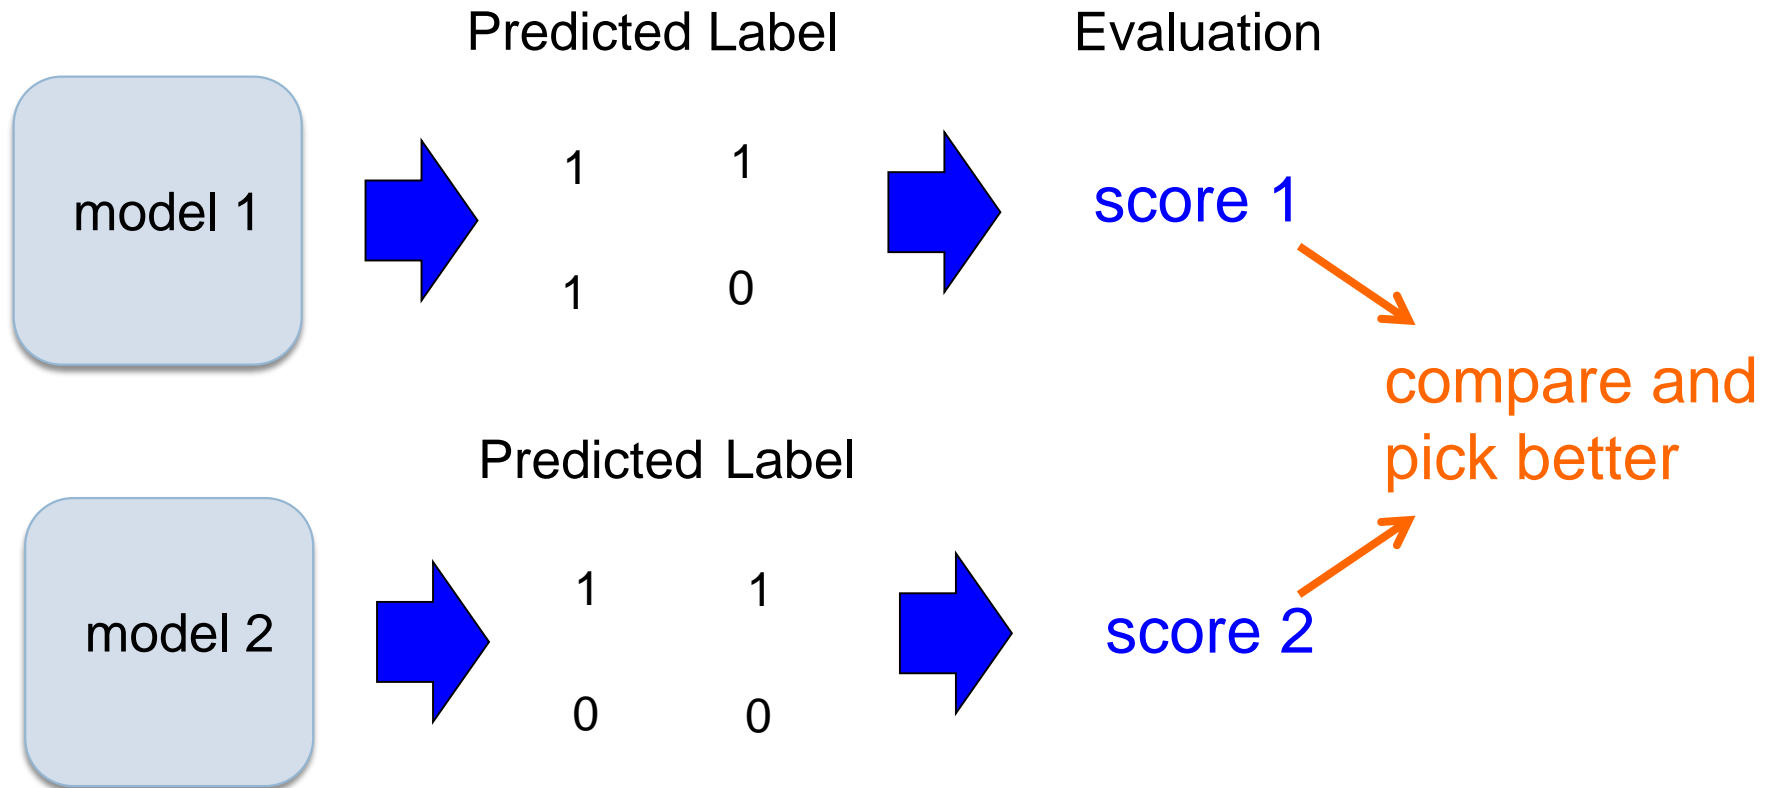
# Idea 1



When would we want to do this type of comparison?



# Idea 1



Any concerns?

# Is model 2 better?

---

Model 1: 85% accuracy

Model 2: 80% accuracy

Model 1: 85.5% accuracy

Model 2: 85.0% accuracy

Model 1: 0% accuracy

Model 2: 100% accuracy

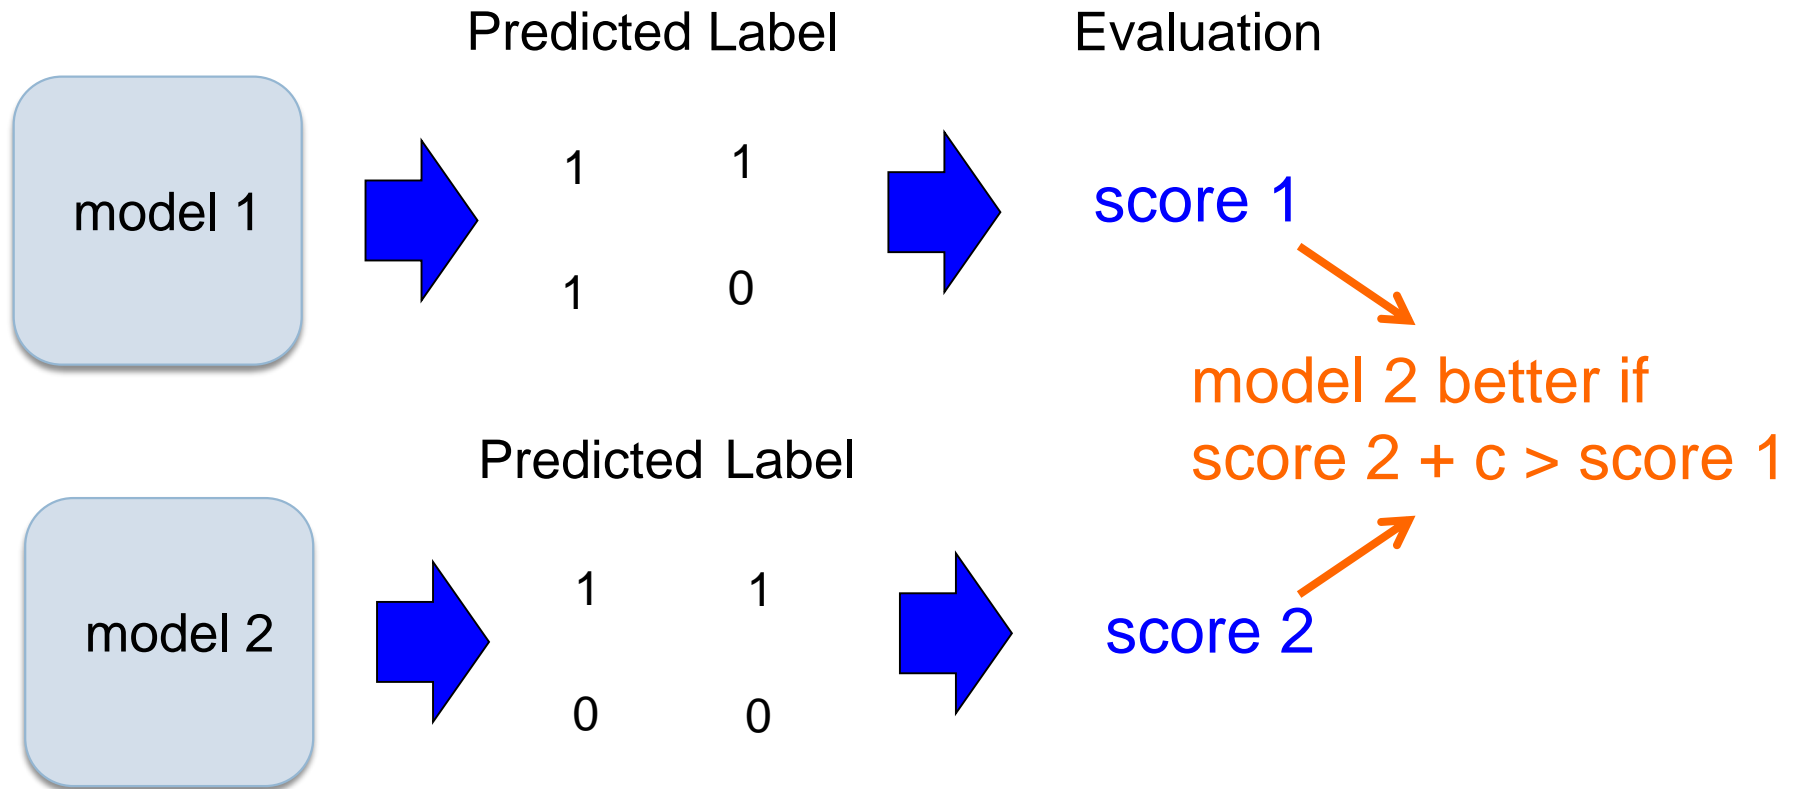
# Comparing scores: significance

Just comparing scores on one data set isn't enough!

We don't just want to know which system is better on *this particular data*, we want to know if model 1 is better than model 2 *in general*

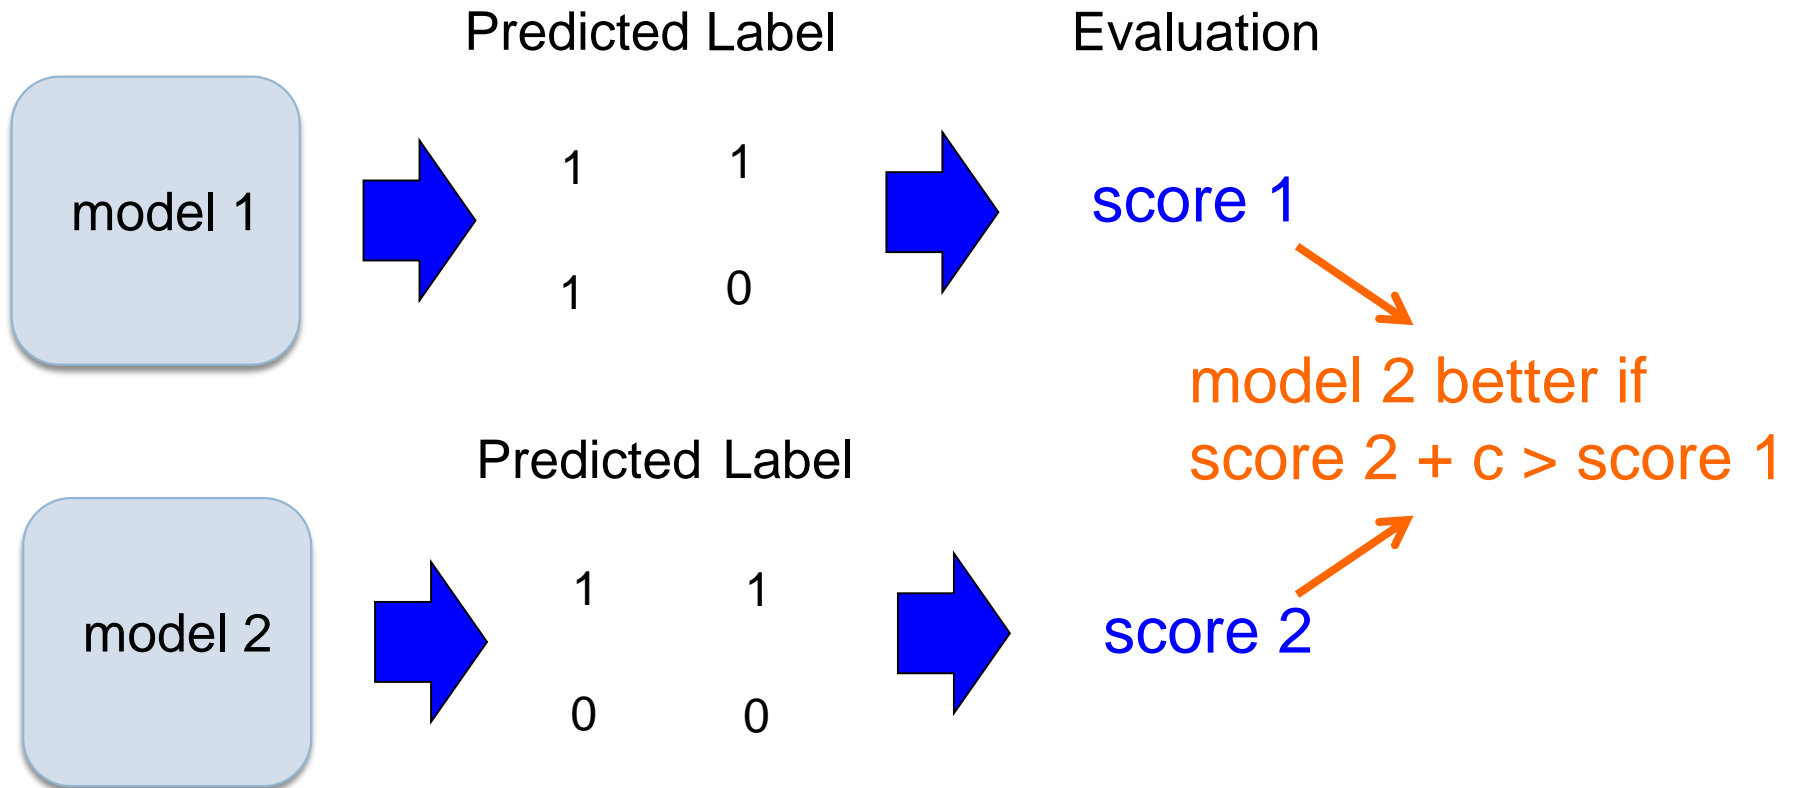
Put another way, we want to be confident that the difference is real and not just due to random chance

# Idea 2



Is this any better?

# Idea 2




**NO!**

**Key:** we don't know the variance of the output

# Variance


Recall that variance (or standard deviation) helped us predict how likely certain events are:



How do we know how variable a model's accuracy is?








# Variance

Recall that variance (or standard deviation) helped us predict how likely certain events are:



We need multiple accuracy scores! Ideas?

# Repeated experimentation

Labeled data	Data	Label	
		0	Training data
		0	
		1	
		1	
		0	
		1	Testing data
		0	







Rather than just splitting once, split multiple times









# Repeated experimentation

Training data







Data Label

	0
	0
	1
	1
	0
	1

Data Label

	0
	0
	1
	1
	0
	1

Data Label

	0
	0
	1
	1
	0
	1

...

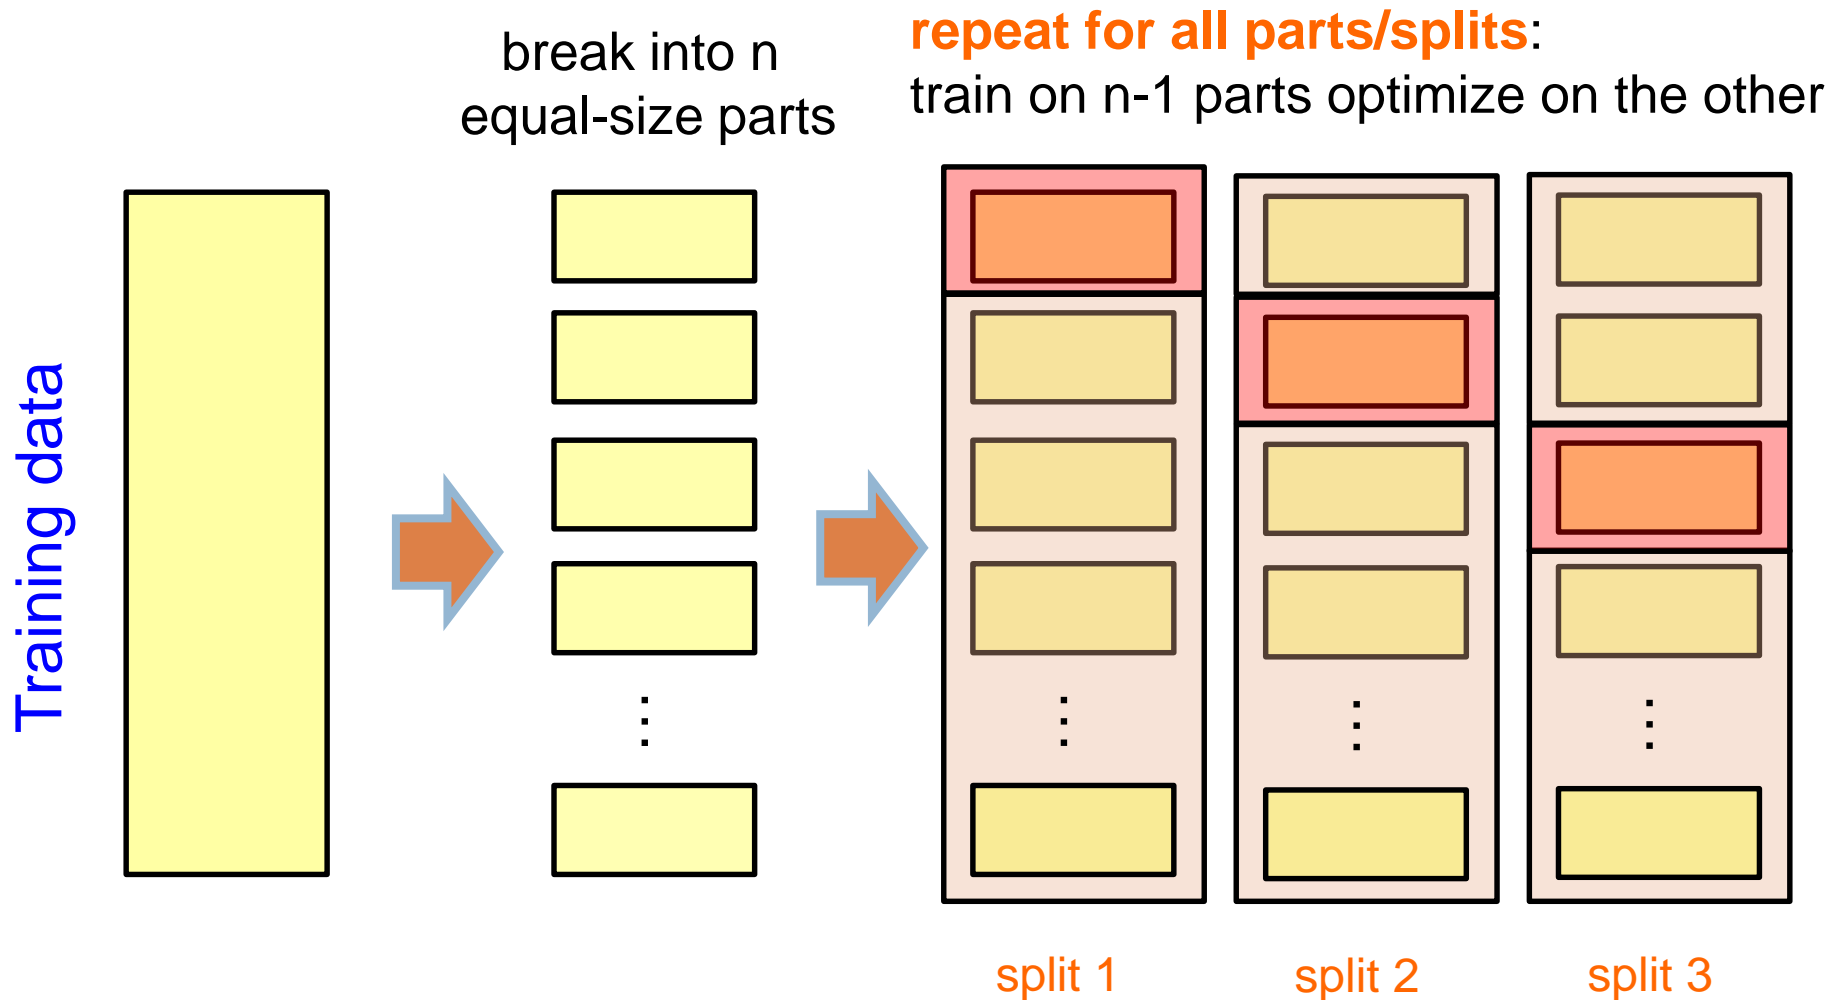


= train




= development

# n-fold cross validation



# n-fold cross validation



# n-fold cross validation

---

better utilization of labeled data

more robust: don't just rely on one test/development set to evaluate the approach (or for optimizing parameters)

multiplies the computational overhead by  $n$  (have to train  $n$  models instead of just one)

10 is the most common choice of  $n$

# Leave-one-out cross validation

---

n-fold cross validation where  $n$  = number of examples

aka “jackknifing”

pros/cons?

when would we use this?

# Leave-one-out cross validation

---

Can be very expensive if training is slow and/or if there are a large number of examples

Useful in domains with limited training data:  
*maximizes the data we can use for training*

# Comparing systems: sample 1

split	model 1	model 2
1	87	88
2	85	84
3	83	84
4	80	79
5	88	89
6	85	85
7	83	81
8	87	86
9	88	89
10	84	85
<b>average:</b>	<b>85</b>	<b>85</b>

Is model 2 better  
than model 1?

# Comparing systems: sample 2

split	model 1	model 2
1	87	87
2	92	88
3	74	79
4	75	86
5	82	84
6	79	87
7	83	81
8	83	92
9	88	81
10	77	85
<b>average:</b>	<b>82</b>	<b>85</b>

Is model 2 better  
than model 1?



# Comparing systems: sample 3

split	model 1	model 2
1	84	87
2	83	86
3	78	82
4	80	86
5	82	84
6	79	87
7	83	84
8	83	86
9	85	83
10	83	85
<b>average:</b>	<b>82</b>	<b>85</b>

Is model 2 better  
than model 1?

# Comparing systems

split	model 1	model 2
1	84	87
2	83	86
3	78	82
4	80	86
5	82	84
6	79	87
7	83	84
8	83	86
9	85	83
10	83	85
average :	82	85

split	model 1	model 2
1	87	87
2	92	88
3	74	79
4	75	86
5	82	84
6	79	87
7	83	81
8	83	92
9	88	81
10	77	85
average :	82	85

What's the difference?

# Comparing systems

split	model 1	model 2
1	84	87
2	83	86
3	78	82
4	80	86
5	82	84
6	79	87
7	83	84
8	83	86
9	85	83
10	83	85
<b>average :</b>	<b>82</b>	<b>85</b>
<b>std dev</b>	<b>2.3</b>	<b>1.7</b>

split	model 1	model 2
1	87	87
2	92	88
3	74	79
4	75	86
5	82	84
6	79	87
7	83	81
8	83	92
9	88	81
10	77	85
<b>average :</b>	<b>82</b>	<b>85</b>
<b>std dev</b>	<b>5.9</b>	<b>3.9</b>

Even though the averages are same, the variance is different!

# Comparing systems: sample 4

split	model 1	model 2
1	80	82
2	84	87
3	89	90
4	78	82
5	90	91
6	81	83
7	80	80
8	88	89
9	76	77
10	86	88
<b>average:</b>	<b>83</b>	<b>85</b>
<b>std dev</b>	<b>4.9</b>	<b>4.7</b>

Is model 2 better  
than model 1?

# Comparing systems: sample 4

split	model 1	model 2	model 2 – model 1
1	80	82	2
2	84	87	3
3	89	90	1
4	78	82	4
5	90	91	1
6	81	83	2
7	80	80	0
8	88	89	1
9	76	77	1
10	86	88	2
<b>average :</b>	<b>83</b>	<b>85</b>	
<b>std dev</b>	<b>4.9</b>	<b>4.7</b>	

Is model 2 better  
than model 1?

# Comparing systems: sample 4

split	model 1	model 2	model 2 – model 1
1	80	82	2
2	84	87	3
3	89	90	1
4	78	82	4
5	90	91	1
6	81	83	2
7	80	80	0
8	88	89	1
9	76	77	1
10	86	88	2
<b>average :</b>	<b>83</b>	<b>85</b>	
<b>std dev</b>	<b>4.9</b>	<b>4.7</b>	

Model 2 is  
ALWAYS better

# Comparing systems: sample 4

split	model 1	model 2	model 2 – model 1
1	80	82	2
2	84	87	3
3	89	90	1
4	78	82	4
5	90	91	1
6	81	83	2
7	80	80	0
8	88	89	1
9	76	77	1
10	86	88	2
<b>average :</b>	<b>83</b>	<b>85</b>	
<b>std dev</b>	<b>4.9</b>	<b>4.7</b>	

How do we decide  
if model 2 is better  
than model 1?

# Statistical tests

## Setup:

- ▣ Assume some default hypothesis about the data that you'd like to *disprove*, called the **null hypothesis**
- ▣ e.g. model 1 and model 2 are not statistically different in performance



## Test:

- ▣ Calculate a test statistic from the data (often assuming something about the data)
- ▣ Based on this statistic, with *some probability* we can **reject the null hypothesis**, that is, show that it does not hold



# t-test

Determines whether two samples come from the same underlying distribution or not



# t-test



Null hypothesis: model 1 and model 2 accuracies are no different, i.e. come from **the same** distribution

Assumptions: there are a number that often aren't completely true, but we're often not too far off

Result: probability that the difference in accuracies is due to random chance (low values are better)

# Calculating t-test

For our setup, we'll do what's called a "pair t-test"

- ▣ The values can be thought of as pairs, where they were calculated under the same conditions
- ▣ In our case, the same train/test split
- ▣ Gives more power than the unpaired t-test (we have more information)

For almost all experiments, we'll do a "two-tailed" version of the t-test

Can calculate by hand or in code, but why reinvent the wheel: use excel or a statistical package

[http://en.wikipedia.org/wiki/Student's\\_t-test](http://en.wikipedia.org/wiki/Student's_t-test)

<http://www.statskingdom.com/160MeanT2pair.html>

<https://www.socscistatistics.com/tests/ttestdependent/Default2.aspx>

# p-value

The result of a statistical test is often a p-value

p-value: the probability that the null hypothesis holds. Specifically, if we re-ran this experiment multiple times (say on different data) what is the probability that we would reject the null hypothesis incorrectly (i.e. the probability we'd be wrong)

Common values to consider “significant”: 0.05 (95% confident), 0.01 (99% confident) and 0.001 (99.9% confident)

# Comparing systems: sample 1

split	model 1	model 2
1	87	88
2	85	84
3	83	84
4	80	79
5	88	89
6	85	85
7	83	81
8	87	86
9	88	89
10	84	85
<b>average:</b>	<b>85</b>	<b>85</b>

Is model 2 better  
than model 1?

They are the same with:  
 $p = 1$

# Comparing systems: sample 2

split	model 1	model 2
1	87	87
2	92	88
3	74	79
4	75	86
5	82	84
6	79	87
7	83	81
8	83	92
9	88	81
10	77	85
<b>average:</b>	<b>82</b>	<b>85</b>

Is model 2 better  
than model 1?

They are the same with:  
 $p = 0.15$

# Comparing systems: sample 3

split	model 1	model 2
1	84	87
2	83	86
3	78	82
4	80	86
5	82	84
6	79	87
7	83	84
8	83	86
9	85	83
10	83	85
average:	82	85

Is model 2 better  
than model 1?

They are the same with:  
 $p = 0.007$

# Comparing systems: sample 4

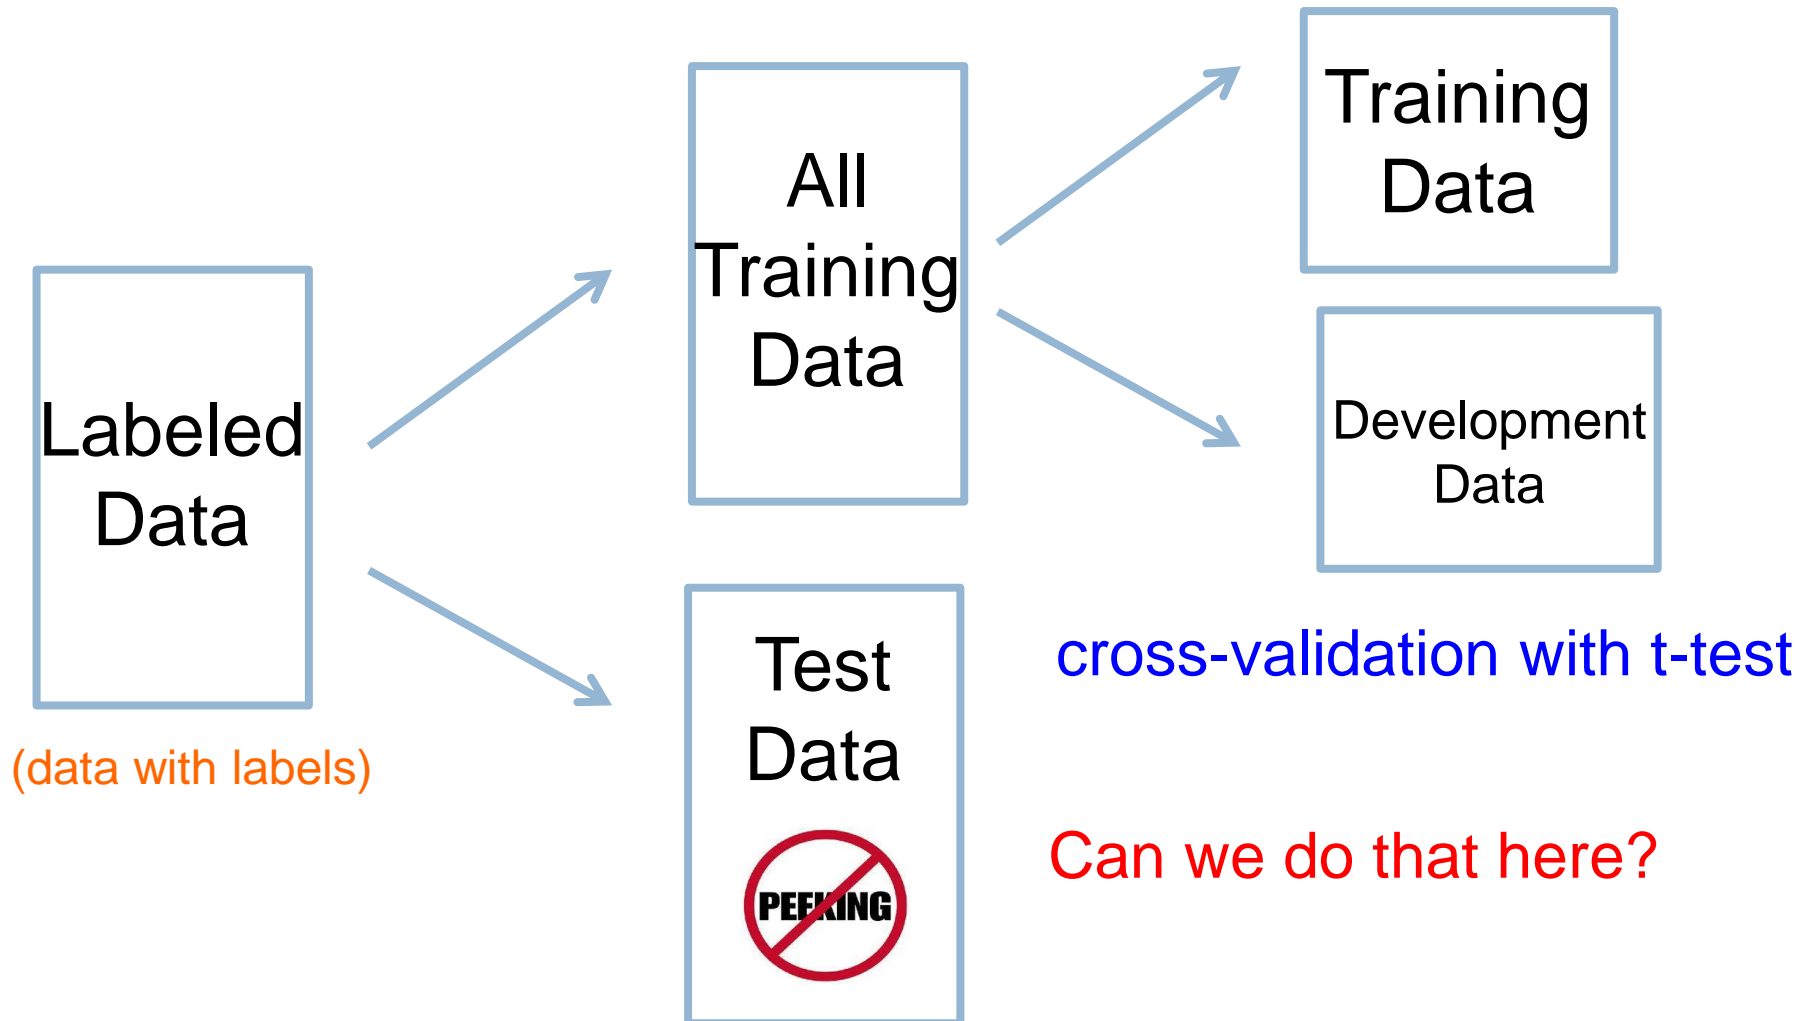
split	model 1	model 2
1	80	82
2	84	87
3	89	90
4	78	82
5	90	91
6	81	83
7	80	80
8	88	89
9	76	77
10	86	88
<b>average:</b>	<b>83</b>	<b>85</b>

Is model 2 better  
than model 1?

They are the same with:  
 $p = 0.001$



# Statistical tests on test data



# Bootstrap resampling

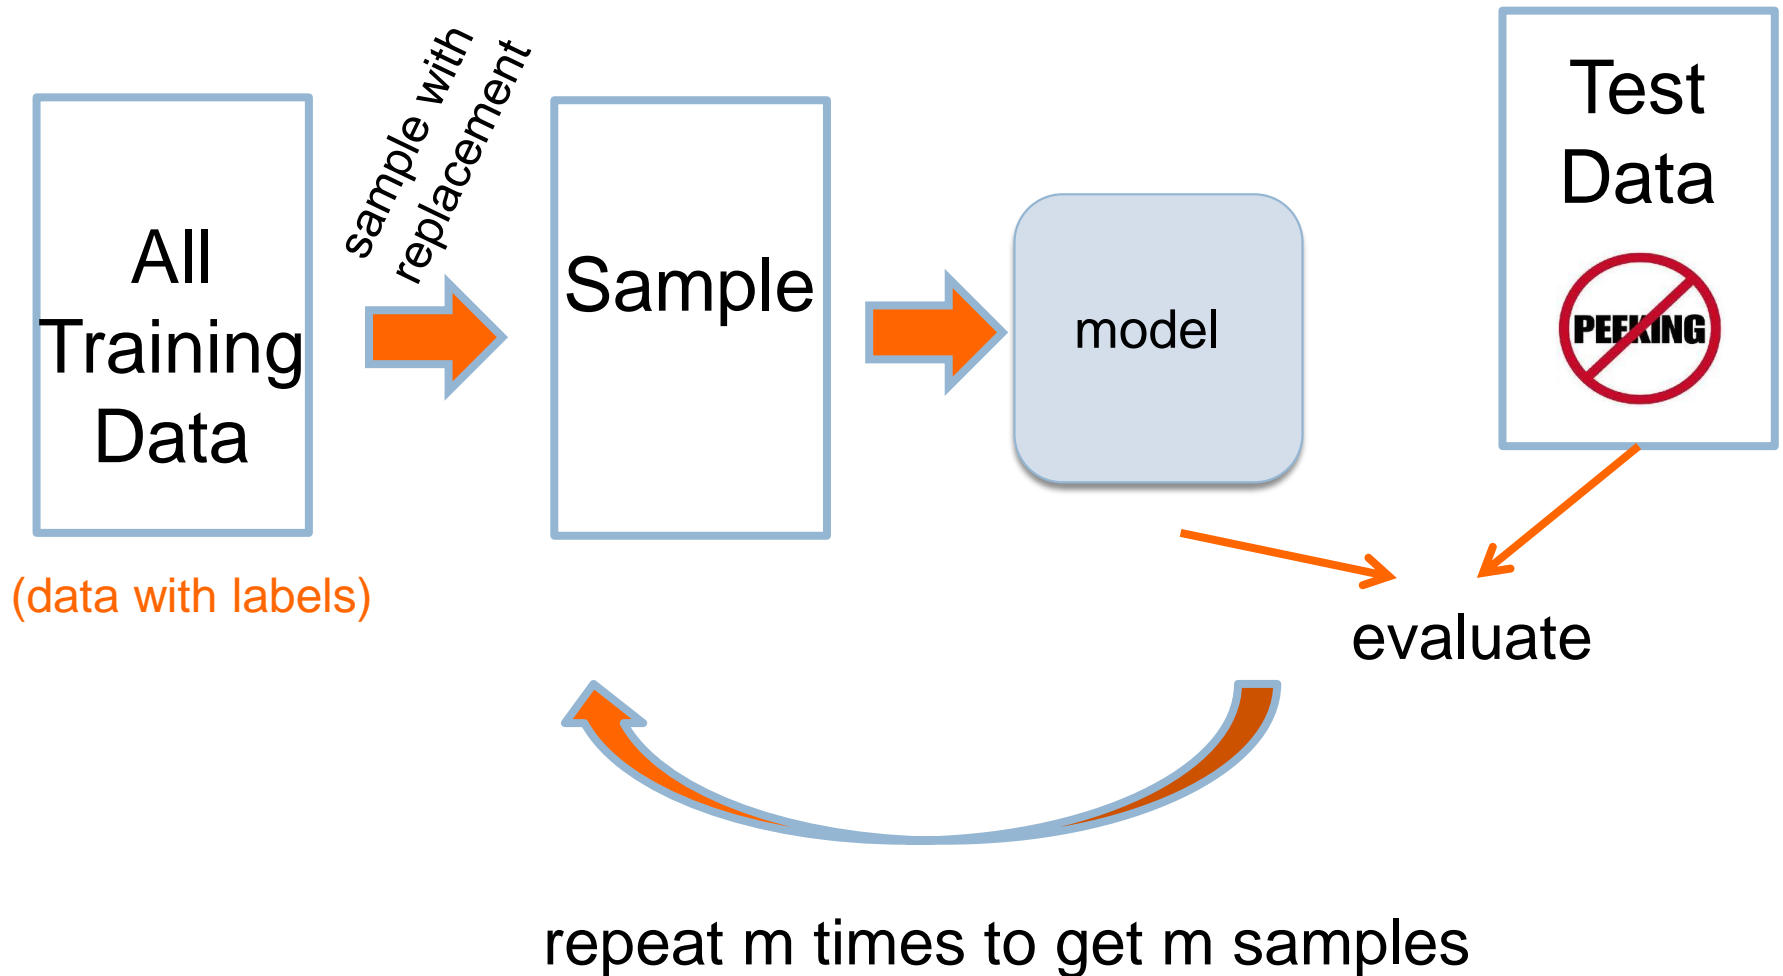
training set  $t$  with  $n$  samples

do  $m$  times:

- sample  $n$  examples **with replacement** from the training set to create a new training set  $t'$
- train model(s) on  $t'$
- calculate performance on test set

calculate t-test (or other statistical test) on the collection of  $m$  results

# Bootstrap resampling



# Experimentation good practices

Never look at your test data!

During development

- ▣ Compare different models/hyperparameters on development data
- ▣ use cross-validation to get more consistent results
- ▣ If you want to be confident with results, use a t-test and look for  $p = 0.05$

For final evaluation, use bootstrap resampling combined with a t-test to compare final approaches

# CSE419 – Artificial Intelligence and Machine Learning 2020

PhD Furkan Gözükar, Toros University

<https://github.com/FurkanGozukara/CSE419-Artificial-Intelligence-and-Machine-Learning-2020>


## Lecture 9 Part 2

### Multiclass

*Based on Asst. Prof. Dr. David Kauchak (Pomona College) Lecture Slides*

# Multiclass classification

## examples



label

Same setup where we have a set of features for each example

apple



orange



apple

Rather than just two labels, now have 3 or more



banana

real-world examples?



banana




pineapple


# Real world multiclass classification



document classification



protein classification




handwriting recognition



face recognition

most real-world applications  
tend to be multiclass



sentiment analysis






autonomous vehicles



emotion recognition

# Multiclass: current classifiers



Any of these work out of the box?  
With small modifications?




# k-Nearest Neighbor (k-NN)

To classify an example  $\mathbf{d}$ :

- ▣ Find  $k$  nearest neighbors of  $\mathbf{d}$
- ▣ Choose as the label the majority label within the  $k$  nearest neighbors

No algorithmic changes!



# Decision Tree learning

Base cases:


1. If all data belong to the same class, pick that label
2. If all the data have the same feature values, pick majority label
3. If we're out of features to examine, pick majority label
4. If the we don't have any data left, pick majority label of *parent*
5. *If some other stopping criteria* exists to avoid overfitting, pick majority label

Otherwise:

- calculate the “score” for each feature if we used it to split the data
- pick the feature with the highest score, partition the data based on that data value and call recursively

No algorithmic changes!


# Perceptron learning



Hard to separate three classes with just one line ☹️

# Black box approach to multiclass


Abstraction: we have a generic binary classifier, how can we use it to solve our new problem



Can we solve our multiclass problem with this?

# Multiclass classification

## examples



label

Same setup where we have a set of features for each example

apple



orange



apple

Rather than just two labels, now have 3 or more



banana




banana



pineapple

# Black box approach to multiclass

Abstraction: we have a generic binary classifier, how can we use it to solve our new problem























Can we solve our multiclass problem with this?


# Approach 1: One vs. all (OVA)

Training: for each label  $L$ , pose as a binary problem

- ▣ all examples with label  $L$  are positive
- ▣ all other examples are negative


		apple vs. not		orange vs. not		banana vs. not	
	apple		+1		-1		-1
	orange		-1		+1		-1
	apple		+1		-1		-1
	banana		-1		-1		+1
	banana		-1		-1		+1

# OVA: linear classifiers (e.g. perceptron)







# OVA: linear classifiers (e.g. perceptron)




# OVA: linear classifiers (e.g. perceptron)




# OVA: linear classifiers (e.g. perceptron)




# OVA: linear classifiers (e.g. perceptron)



# OVA: linear classifiers (e.g. perceptron)



# OVA: linear classifiers (e.g. perceptron)




# OVA: classify

## Classify:


- ▣ If classifier doesn't provide confidence (this is rare) and there is ambiguity, pick one of the ones in conflict
- ▣ Otherwise:
  - pick the most confident positive
  - if none vote positive, pick *least* confident negative

# OVA: linear classifiers (e.g. perceptron)





# OVA: linear classifiers (e.g. perceptron)



# OVA: classify, perceptron

## Classify:

- ▣ If classifier doesn't provide confidence (this is rare) and there is ambiguity, pick majority in conflict
- ▣ Otherwise:
  - pick the most **confident** positive
  - if none vote positive, pick *least* confident negative

How do we calculate this for the perceptron?

# OVA: classify, perceptron

## Classify:

- ▣ If classifier doesn't provide confidence (this is rare) and there is ambiguity, pick majority in conflict
- ▣ Otherwise:
  - pick the most **confident** positive
  - if none vote positive, pick *least* confident negative

$$\text{prediction} = b + \sum_{i=1}^n w_i f_i$$

Distance from the hyperplane

# Approach 2: All vs. all (AVA)

Training:

For each pair of labels, train a classifier to distinguish between them

for  $i = 1$  to number of labels:






for  $k = i+1$  to number of labels:

train a classifier to distinguish between  $label_j$  and  $label_k$ :

- create a dataset with all examples *with  $label_j$*  labeled positive and all examples with  $label_k$  labeled negative

- train classifier on this subset of the data

# AVA training visualized

	<b>apple</b>
	<b>orange</b>
	<b>apple</b>
	<b>banana</b>
	<b>banana</b>

## apple vs orange



## orange vs banana



## apple vs banana



# AVA classify

apple vs orange



+1



+1



-1

apple vs banana



+1



+1



-1



-1

orange vs banana



+1



-1



-1



What class?

# AVA classify

apple vs orange



+1



+1

orange



-1

apple vs banana



+1



+1

apple



-1



-1

orange vs banana



+1



-1

orange



-1



orange

In general?

# AVA classify

To classify example  $e$ , classify with each classifier  $f_{jk}$

We have a few options to choose the final class:

- Take a majority vote
- Take a weighted vote based on confidence
  - $y = f_{jk}(e)$
  - $\text{score}_j += y$  **How does this work?**
  - $\text{score}_k -= y$

*Here we're assuming that  $y$  encompasses both the prediction (+1, -1) and the confidence, i.e.  $y = \text{prediction} * \text{confidence}$ .*



# AVA classify

Take a weighted vote based on confidence

- $y = f_{jk}(e)$
- $\text{score}_j += y$
- $\text{score}_k -= y$

If  $y$  is positive, classifier thought it was of type  $j$ :

- raise the score for  $j$
- lower the score for  $k$

if  $y$  is negative, classifier thought it was of type  $k$ :

- lower the score for  $j$
- raise the score for  $k$

# OVA vs. AVA

---

Train/classify runtime?

Error? Assume each binary classifier makes an error with probability  $\epsilon$

# OVA vs. AVA

Train time:

AVA learns more classifiers, however, they're trained on much smaller data this tends to make it faster if the labels are equally balanced


Test time:

AVA has more classifiers

Error (see the book for more justification):

- AVA trains on more balanced data sets
- AVA tests with more classifiers and therefore has more chances for errors
- Theoretically:
  - OVA:  $\varepsilon$  (number of labels - 1)
  - AVA:  $2 \varepsilon$  (number of labels - 1)

# Approach 3: Divide and conquer



Pros/cons vs. AVA?

# Multiclass summary







---

If using a binary classifier, the most common thing to do is OVA

Otherwise, use a classifier that allows for multiple labels:

- ▣ DT and k-NN work reasonably well
- ▣ We'll see a few more in the coming weeks that will often work better

# Multiclass evaluation

	label	prediction
	apple	orange
	orange	orange
	apple	apple
	banana	pineapple
	banana	banana
	pineapple	pineapple

How should we evaluate?

# Multiclass evaluation



label

apple

prediction

orange



orange

orange



apple

apple

Accuracy: 4/6



banana

pineapple



banana

banana



pineapple

pineapple

# Multiclass evaluation imbalanced data



label

prediction

apple

orange

...



apple

apple

Any problems?



banana

pineapple

Data imbalance!



banana

banana



pineapple

pineapple



# Macroaveraging vs. microaveraging

**microaveraging**: average over examples (this is the “normal” way of calculating)

**macroaveraging**: calculate evaluation score (e.g. accuracy) for each label, then average over labels

What effect does this have?  
Why include it?







# Macroaveraging vs. microaveraging

**microaveraging**: average over examples (this is the “normal” way of calculating)

**macroaveraging**: calculate evaluation score (e.g. accuracy) for each label, then average over labels

- Puts more weight/emphasis on rarer labels
- Allows another dimension of analysis







# Macroaveraging vs. microaveraging

	label	prediction
	apple	orange
	orange	orange
	apple	apple
	banana	pineapple
	banana	banana
	pineapple	pineapple

**microaveraging:**  
average over examples

**macroaveraging:**  
calculate evaluation  
score (e.g. accuracy)  
for each label, then  
average over labels

# Macroaveraging vs. microaveraging

	label	prediction
	apple	orange
	orange	orange
	apple	apple
	banana	pineapple
	banana	banana
	pineapple	pineapple

microaveraging: 4/6

macroaveraging:

apple = 1/2

orange = 1/1

banana = 1/2

pineapple = 1/1

total =  $(1/2 + 1 + 1/2 + 1)/4$

= 3/4


# Confusion matrix

entry  $(i, j)$  represents the number of examples with label  $i$  that were predicted to have label  $j$

another way to understand both the data and the classifier

	Classic	Country	Disco	Hiphop	Jazz	Rock
Classic	86	2	0	4	18	1
Country	1	57	5	1	12	13
Disco	0	6	55	4	0	5
Hiphop	0	15	28	90	4	18
Jazz	7	1	0	0	37	12
Rock	6	19	11	0	27	48

# Confusion matrix



BLAST classification of proteins in 850  
superfamilies